

To start with a simple example, let's say that your goal is to build a logistic regression model in Python in order to determine whether candidates would get admitted to a prestigious university.

You can build a logistic regression in Python, where:

The 3 independent variables are the GMAT score, GPA and Years of work experience

```
In [30]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn import metrics
import seaborn as sn
import matplotlib.pyplot as plt
```

```
In [2]: candidates = {'gmat': [780,750,690,710,680,730,690,720,740,690,610,690,710,680,770,6  
    'gpa': [4,3.9,3.3,3.7,3.9,3.7,2.3,3.3,3.3,1.7,2.7,3.7,3.7,3.3,3.3,3,2.  
    'work_experience': [3,4,3,5,4,6,1,4,5,1,3,5,6,4,3,1,4,6,2,3,2,1,4,1,2,  
    'admitted': [1,1,0,1,0,1,0,1,1,0,0,1,1,0,1,0,0,1,0,0,1,0,0,0,0,1,1,0,1  
    }  
  
df = pd.DataFrame(candidates,columns= ['gmat', 'gpa','work_experience','admitted'])  
print (df)
```

	gmat	gpa	work_experience	admitted
0	780	4.0	3	1
1	750	3.9	4	1
2	690	3.3	3	0
3	710	3.7	5	1
4	680	3.9	4	0
5	730	3.7	6	1
6	690	2.3	1	0
7	720	3.3	4	1
8	740	3.3	5	1
9	690	1.7	1	0
10	610	2.7	3	0
11	690	3.7	5	1
12	710	3.7	6	1
13	680	3.3	4	0
14	770	3.3	3	1
15	610	3.0	1	0
16	580	2.7	4	0
17	650	3.7	6	1
18	540	2.7	2	0
19	590	2.3	3	0
20	620	3.3	2	1
21	600	2.0	1	0

22	550	2.3	4	0
23	550	2.7	1	0
24	570	3.0	2	0
25	670	3.3	6	1
26	660	3.7	4	1
27	580	2.3	2	0
28	650	3.7	6	1
29	660	3.3	5	1
30	640	3.0	1	0
31	620	2.7	2	0
32	660	4.0	4	1
33	660	3.3	6	1
34	680	3.3	5	1
35	650	2.3	1	0
36	670	2.7	2	0
37	580	3.3	1	0
38	590	1.7	4	0
39	690	3.7	5	1

In [3]: `df.describe()`

Out[3]:

	gmat	gpa	work_experience	admitted
count	40.000000	40.000000	40.000000	40.000000
mean	654.000000	3.095000	3.425000	0.475000
std	61.427464	0.631218	1.737778	0.505736
min	540.000000	1.700000	1.000000	0.000000
25%	607.500000	2.700000	2.000000	0.000000
50%	660.000000	3.300000	4.000000	0.000000
75%	690.000000	3.700000	5.000000	1.000000
max	780.000000	4.000000	6.000000	1.000000

In [5]: `print(df.shape)`

(40, 4)

Step 4: Create the dependent & independent variables for logistic regression

In [6]:

```
x = df[['gmat', 'gpa', 'work_experience']]
y = df['admitted']
```

Apply 'train_test_split'. For example, you can set the test size to 0.25, and therefore the model testing will be based on 25% of the dataset, while the model training will be based on 75% of the dataset

In [7]: `X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.25,random_state=0)`

In [8]: `print(X_train.shape)`

(30, 3)


Apply the logistic regression

```
In [9]: logistic_regression= LogisticRegression()
logistic_regression.fit(X_train,y_train)
y_pred=logistic_regression.predict(X_test)
```

Creating the Confusion Matrix

```
In [10]: confusion_matrix = pd.crosstab(y_test, y_pred, rownames=['Actual'], colnames=['Predicted'],
sn.heatmap(confusion_matrix, annot=True)
```

```
Out[10]: <AxesSubplot:xlabel='Predicted', ylabel='Actual'>
```



print the Accuracy and plot the Confusion Matrix

```
In [11]: print('Accuracy: ',metrics.accuracy_score(y_test, y_pred))
plt.show()
```

Accuracy: 0.8

As can be observed from the matrix:

TP = True Positives = 4 TN = True Negatives = 4 FP = False Positives = 1 FN = False Negatives = 1

Accuracy = (TP+TN)/Total = (4+4)/10 = 0.8

Diving Deeper into the Results

```
In [14]: print (X_test)

##Recall that our original dataset (from step 1) had 40 observations. Since we set t
```

	gmat	gpa	work_experience
22	550	2.3	4
20	620	3.3	2
25	670	3.3	6
4	680	3.9	4
10	610	2.7	3
15	610	3.0	1
28	650	3.7	6

11	690	3.7	5
18	540	2.7	2
29	660	3.3	5

In [15]: `#The prediction was also made for those 10 records (where 1 = admitted, while 0 = re`
`print (y_pred)`

`[0 0 1 1 0 0 1 1 0 1]`

In the actual dataset (from step-1), you'll see that for the test data, we got the correct results 8 out of 10 times

Index	gmat	gpa	work_experience	admitted - actual results	admitted - predicted results	Matching
22	550	2.3	4	0	0	TRUE
20	620	3.3	2	1	0	FALSE
25	670	3.3	6	1	1	TRUE
4	680	3.9	4	0	1	FALSE
10	610	2.7	3	0	0	TRUE
15	610	3	1	0	0	TRUE
28	650	3.7	6	1	1	TRUE
11	690	3.7	5	1	1	TRUE
18	540	2.7	2	0	0	TRUE
29	660	3.3	5	1	1	TRUE

Checking the Prediction for a New Set of Data

Let's say that you have a new set of data, with 5 new candidates

Goal is to use the existing logistic regression model to predict whether the new candidates will get admitted

In [16]: `##Creating the new candidates dataframe`
`new_candidates = {'gmat': [590,740,680,610,710],`
`'gpa': [2,3.7,3.3,2.3,3],`
`'work_experience': [3,4,6,1,5]`
`}`
`df2 = pd.DataFrame(new_candidates,columns= ['gmat', 'gpa','work_experience'])`

In [17]: `df2.describe()`

Out[17]:

	gmat	gpa	work_experience
count	5.000000	5.00000	5.000000
mean	666.000000	2.86000	3.800000
std	64.265076	0.70214	1.923538
min	590.000000	2.00000	1.000000
25%	610.000000	2.30000	3.000000
50%	680.000000	3.00000	4.000000

	gmat	gpa	work_experience
75%	710.000000	3.30000	5.000000
max	740.000000	3.70000	6.000000

In [18]: `y_pred=logistic_regression.predict(df2)`

In [19]: `print (df2)`

```

   gmat  gpa  work_experience
0   590  2.0                3
1   740  3.7                4
2   680  3.3                6
3   610  2.3                1
4   710  3.0                5

```

In [20]: `print (y_pred)`

```
[0 1 1 0 1]
```

The first and fourth candidates are not expected to be admitted, while the other candidates are expected to be admitted

Calculating ROC Curve

In [23]: `import sklearn.metrics as metrics`

```

# calculate the fpr and tpr for all thresholds of the classification
probs = logistic_regression.predict_proba(X_test)
preds = probs[:,1]
fpr, tpr, threshold = metrics.roc_curve(y_test, preds)
roc_auc = metrics.auc(fpr, tpr)

```

Difference between "The predict() method" and "The predict_proba() method"

The predict() method

All supervised estimators in scikit-learn implement the predict() method that can be executed on a trained model in order to predict the actual label (or class) over a new set of data.

The method accepts a single argument that corresponds to the data over which the predictions will be made and it returns an array containing the predicted label for each data point.

In [25]: `predictions = logistic_regression.predict(X_test)`
`predictions`

Out[25]: `array([0, 0, 1, 1, 0, 0, 1, 1, 0, 1], dtype=int64)`

The predict_proba() method

In the context of classification tasks, some sklearn estimators also implement the predict_proba

method that returns the class probabilities for each data point.


The method accepts a single argument that corresponds to the data over which the probabilities will be computed and returns an array of lists containing the class probabilities for the input data points.

```
In [27]: predictions = logistic_regression.predict_proba(X_test)
print(predictions)
```

```
[[9.91609541e-01 8.39045850e-03]
 [9.82686248e-01 1.73137518e-02]
 [4.82945926e-02 9.51705407e-01]
 [2.31228082e-01 7.68771918e-01]
 [9.72097119e-01 2.79028812e-02]
 [9.97314934e-01 2.68506615e-03]
 [7.39280072e-02 9.26071993e-01]
 [6.18376418e-02 9.38162358e-01]
 [9.99411602e-01 5.88398313e-04]
 [2.11035600e-01 7.88964400e-01]]
```

```
In [28]: # method using: plt

import matplotlib.pyplot as plt
plt.title('Receiver Operating Characteristic')
plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



```
In [29]: y_pred_proba = logistic_regression.predict_proba(X_test)[:,:1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred_proba)
auc = metrics.roc_auc_score(y_test, y_pred_proba)
plt.plot(fpr,tpr,label="data 1, auc="+str(auc))
plt.legend(loc=4)
plt.show()
```

